CSCI 446 Artificial Intelligence Project 1 Final Report

Roy Smart   Nevin Leh   Brian Marsh

September 25, 2016

1 Introduction

The *Graph Coloring Problem* (GCP) is the problem of attempting to color a set of interconnected vertices, using a limited set of colors, such that no vertex has the same color as its neighbors. This problem is best visualized as the map coloring problem. As an example of the map coloring problem, consider the problem of assigning one of four colors to every country in the world, such that no two adjacent countries have the same color, as in Figure 1. It can be shown that the map coloring problem reduces to the graph coloring problem[2]
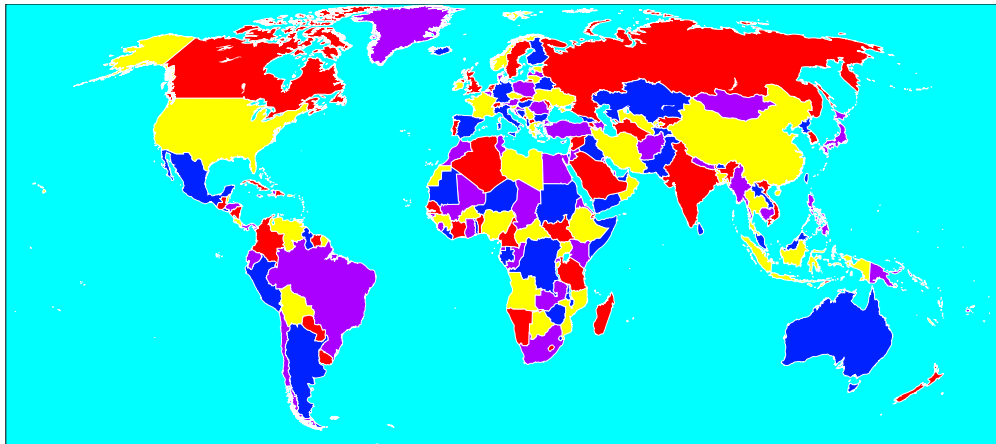
Figure 1: *Map of the world satisfying the map coloring problem [1].*

if we represent the countires as the vertices of the graph, and the borders between countries as the edges of the graph, shown in Figure 2. This configuration produces a *planar* graph, a graph with no edge intersections. For this project, we are asked to solve the GCP problem for planar graphs only. Since the problems are

Figure 2: *Diagram describing how the map coloring problem can be transformed into the graph coloring problem [3].*

guaranteed to be planar, we will represent our graphs using polygon maps generated by our program.

To solve the GCP we employed five different algorithms: Minimum Conflicts, Simple Backtracking, Backtracking with Forward Checking, Backtracking with Constraint Propagation (MAC), and Local Search

using a Genetic Algorithm. To evaluate these algorithms, we built a problem generating program that can produce a random set of planar graphs. Using the problem generator, we calculated a set of graphs between the sizes {10, 20, 30,...,100} and called the five graph coloring algorithms to solve the GCP. We measured GCP algorithm performance by measuring the number of read/write operations, the time required to find a solution, and the number of function calls for each algorithm. Using these metrics, we predict that the Minimum Conflicts algorithm will be the fastest, based off its performance on the eight-queens problem[2].

The code for this project is implemented in `C++` and depends on the programs `cairo`, `gnuplot`, `ffmpeg`, and LaTeX for graph output, performance output, video output, and documentation respectively.

2 PROBLEM GENERATION

We created a function that could produce many examples of planar graphs from a set of randomly scattered points. For an arbitrary set of points, there is no unique planar graph that can be constructed. To solve this issue, the problem statement has provided a algorithm for calculating an planar graph.

> Select some point X at random and connect X by a straight line to the nearest point Y such that X is not already connected to Y and line crosses no other line. Repeat the previous step until no more connections are possible.

For our implementation of the above algorithm, we started by creating a complete graph (where each vertex is connected to every other vertex). Each vertex was associated with a list of edges, sorted by length. We then selected a point at random and inspected the first unchecked edge E. If E did not cross any of the accepted edges, it is added to the list of accepted edges and marked as checked. This process was repeated until every edge was marked as checked.

The most challenging part of the implementation described above was determining if two line segments intersected. An algorithm to determine if two edges (or line segments) cross has been outlined by LaMothe [4] and described in detail here. Let's start by defining two line segments

$$\begin{aligned} A &= \{\mathbf{a},\ \mathbf{a'}\} \\ B &= \{\mathbf{b},\ \mathbf{b'}\} \end{aligned} \tag{1}$$

where $\mathbf{a}$, $\mathbf{a'}$, $\mathbf{b}$, and $\mathbf{b'}$ are vectors from the origin to the ends of the line segments. Next, compute the direction vectors for each line segment

$$\begin{aligned} \boldsymbol{\alpha} &= \mathbf{a'} - \mathbf{a} \\ \boldsymbol{\beta} &= \mathbf{b'} - \mathbf{b} \end{aligned} \tag{2}$$

Now, the trick to solving this problem is to parameterize the line segment using the direction vector and a parameter t_i in the domain $[0, 1]$.

$$\begin{aligned} \mathbf{p} &= \mathbf{a} + \boldsymbol{\alpha} t_a, \quad t_a \in [0,\ 1] \\ \mathbf{q} &= \mathbf{b} + \boldsymbol{\beta} t_b, \quad t_b \in [0,\ 1] \end{aligned} \tag{3}$$

From here, the solution is obvious. To find the point of intersection, we set $\mathbf{p} = \mathbf{q}$ and solve for the values of t_a and t_b. Then, if the solution is outside the domain of t_i, the line segments do not intersect. Thus,

$$\Rightarrow \begin{cases} p_x = q_x \\ p_y = q_y \end{cases} \tag{4}$$

$$\Rightarrow \begin{cases} a_x + \alpha_x t_a = b_x + \beta_x t_b \\ a_y + \alpha_y t_a = b_y + \beta_y t_b. \end{cases} \tag{5}$$

Equation 5 is a system of two equations and two unknowns, solving for t_a and t_b gives

$$t_a = \frac{(-a_y + b_y)\beta_x + (a_x - b_x)\beta_y}{\alpha_y\beta_x - \alpha_x\beta_y}$$
$$t_b = \frac{(a_y - b_y)\alpha_x + (-a_x + b_x)\alpha_y}{-\alpha_y\beta_x + \alpha_x\beta_y}. \tag{6}$$

Finally, we can determine whether A and B intersect if $0 < t_a < 1$ and $0 < t_b < 1$ evaluates to true. This method may be more inefficient than other algorithms, such as the example outlined by Cormen[5] due to the division operation in Equation 6.

Using the expression above to determine line intersections, we were able to implement a problem generator that was reasonably efficient, generating approximately 1×10^3 examples per second. This efficiency was important, because it made testing and debugging the GCP algorithms much faster. Example output from the problem generator is shown in Figure 3a.

(a) *Planar graph with 100 vertices.*

(b) *Polygon map corresponding to (a).*

Figure 3: *Visualization of the output of the problem generator. First a planar graph was created (a), and then a corresponding polygon map was found (b).*

We have also implemented a way to produce a map of regions that is analogous to the graphs created by the problem generator, shown in Figure 3b. Since the graph coloring problem does not uniquely define the map coloring problem, we had to invent a metric that produces appropriate graphs. To define our metric, consider that if a polygon G_i corresponding to each point P_i were constructed out of the midpoints $M_{i,m}$ of the associated edges $E_{i,m}$ of P_i, then an adjacent polygon G_j would only touch at the corners, that is, our map would be full of gaps. Therefore, to fix the problem we also define the centroid $C_{i,n}$ of each triangle $T_{i,n}$ formed by the two adjacent edges $E_{i,m}$ and $E_{i,m+1}$. We then define the final polygon for P_i using the points $M_{i,m}$ and $C_{i,n}$. This procedure is unnecessary in terms of the assignment, but is helpful for debugging GCP algorithms, demonstrating example output, and for aesthetic pleasure.

3   Minimum Conflicts

4   Simple Backtracking

4.1   Description

Backtracking is a method of solving constraint satisfaction problems through an "educated" exhaustive search. It works by building a partial solution one component at a time and then testing the solution to check if the solution still has a chance of success [6]. If the solution cannot possibly succeed, the algorithm *backtracks* by moving back one component, modifying the component to form a new partial solution and then continuing. By checking partial solutions, backtracking can eliminate large portions of the solution-space, thus increasing the efficiency of the search.

     In the graph coloring problem, backtracking assigns colors to one vertex at a time and checks to see if it has the same value as any of its neighbors. If none of the possible colors work for the current vertex, it backs up and tries a different color on a previously assigned vertex. If a color has no conflicts, backtracking continues to the next vertex in the sequence.
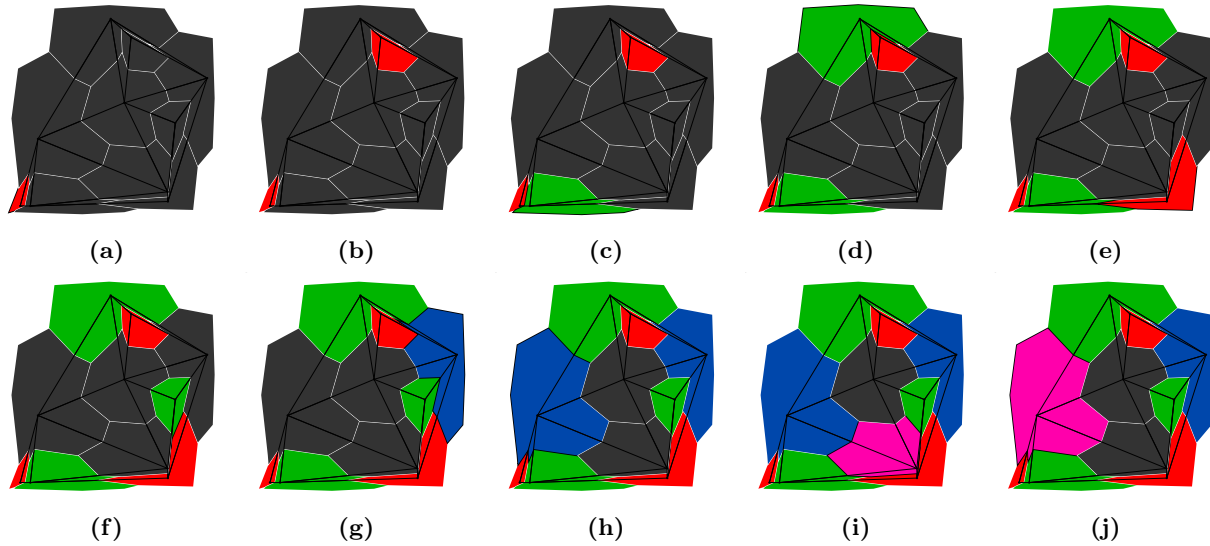
Figure 4: *A map coloring example using simple backtracking on a graph with 10 vertices. From this diagram we will define a label for each vertex by the panel during which it is assigned, for example vertex (a) is colored red in the first panel, vertex (c) is colored green in the third panel, etc. Note we will define vertex (j) as the last uncolored node, since panel (j) changes the value on vertex (h).*

     To illustrate the process of backtracking consider the sequence of images in Figure 4. This sequence follows the backtracking algorithm assigning the first acceptable color (in the order red, green, pink) to each vertex until the first backtracking event. The colors assigned to the vertices (a) through (g) seem reasonable until blue is assigned to vertex (h). Then we can see that the only acceptable color for vertex (i) is pink and even worse, vertex (j) has cannot be assigned a color without being in conflict with it's neighbors. Therefore the algorithm is forced to backtrack to vertex (h) and try pink (which also will not work).

     The example above seeks to describe the major issue with backtracking: it can get stuck in situations where it can take an inordinate amount of time to backtrack far enough to fix the problem. In panel (j), after the backtracking operation, we can see that the algorithm did not retreat far enough, changing vertex (g) to pink still results in a conflict in vertex (j). The algorithm needs to back up to vertex (e) and change it to blue to resolve the conflict on vertex (j) and solve the graph. We will see in the following sections on forward checking and constraint propagation if we can mitigate this problem.

To implement simple backtracking we followed the recursive algorithm outlined in Russel and Norvig Section 6.3 [2]. Russel's BACKTRACK loops through each value in the domain, and if the value is consistent with the constraints it recursively calls BACKTRACK. The function returns `true` when every variable has been assigned, and `false` if it cannot satisfy the constraints. The recursive function is advantageous, as it saves the state of the stack before each recursive call, making it trivial to restore the state of the solution after a failed recursive call to backtracking.

The crux of the backtracking procedure in the graph coloring problem is checking whether a coloring is consistent with constraints. From profiling our code, it was obvious that a fast constraint checking function was critical to the ability of backtracking to solve graph coloring problems up to $N = 100$ vertices in a reasonable amount of time. This facilitated a change of design from an object-oriented approach, to an array-based approach.

4.2 Experimental Approach

To evaluate backtracking's performance on the graph coloring problem, we ran the algorithm on ten graph sizes $N = (10, 20, ..., 100)$ with 20 trials for each value, for a total of 200 backtracking experiments. From each experiment we recorded the number of vertex reads, vertex writes, and number of recursive calls. Time constraints on this project demanded that we set a limit on the number of calls to backtracking, therefore we also recorded if this limit was reached for each experiment.

We expect that the number of recursive calls and the number of vertex writes and the number of vertex reads to be correlated by some constant factor. This is because there is a maximum of two writes and N reads for each recursive call. Therefore, we can use this information to verify that the algorithm is working properly.

4.3 Results

Figure 5: *Logarithmic plot describing the performance of the backtracking algorithm vs. the number of vertices. The shaded regions represent the minimum and maximum values of each quantity.*

5 BACKTRACKING WITH FORWARD CHECKING

5.1 Description

Backtracking with forward checking incorporates the usual simple backtracking algorithm described in the preceding section, while incorporating an ability to look into the future to see if a particular assignment would create an inconsistency. Forward checking achieves this by trying to make an inconsistency occur as fast as possible by checking if all current assignments are consistent with each other [7].

To accomplish forward checking in the graph coloring problem, we need to modify our representation of the graph to allow superpositions of colors and initialize the graph such that all colors are possible for each vertex. Next, for each call to backtracking, we assign a color to the next vertex, and delete that color from the vertex's neighbors list of possible colors. If any vertex has its list of possible colors completely eliminated, the algorithm backtracks and changes a color further up the tree.[2].

Figure 6: *The map coloring example from Figure 4 using backtracking with forward checking. Note that the intermediate colors represent additive mixings of the four pure colors (red, green, blue, and pink), for example purple (the color of vertex (h) during panel (c)) is a combination of blue and pink, magenta is a combination of red and blue, etc.*

To visualize how backtracking with forward checking is different from simple backtracking, we will consider again the map coloring example from the preceding section. To represent the superposition of colors required by this algorithm, we have formed a color map that represents additive mixings of the four pure colors. In this scheme, the color with all colors activated is mauve [RGBP]. The example starts with the graph initialized to mauve. Then vertex (a) is set to red, deleting red from its vertices (h) and (c), forming gray [GBP]. After a similar process for vertex (b), vertex (c) is assigned green, deleting green from vertices (e) and (i) to form magenta [RBP] and from vertex (h) to form purple [RP].

This process produces similar results to simple backtracking until vertex (g) is colored blue, it then deletes blue from vertex (j), leaving pink. Next, vertex (h) is chosen to be blue, leaving vertex (i) with pink as it's only possible color. The algorithm then selects the only choice for vertex (i) (pink), leaving vertex (j) with the null set. Since the null set is not allowed, the algorithm backtracks to vertex (g) and changes its color to pink.

In this example, forward checking had no advantages over simple backtracking, as it discovered the inconsistency in vertex (j) in approximately the same number of steps. However, forward checking offers immense advantages for large graphs, where it can check if a color is consistent with its neighbors much sooner than simple backtracking, leading to much faster solutions.

5.2   Experimental Approach

Our experimental approach for this algorithm is essentially the same as the approach for simple backtracking described in Section 4.2. We will still be measuring the number of vertex reads, vertex writes, and the number of recursive calls. However, since the forward checking algorithm must carry out more decisions for each recursive call, we expect there to be a larger amount of vertex reads and vertex writes for each recursive call. We will see in the next section if this additional overhead was worth it compared to simple backtracking.

5.3   Results

Figure 7: *Logarithmic plot describing the performance of the backtracking algorithm vs. the number of vertices. The shaded regions represent the minimum and maximum values of each quantity.*
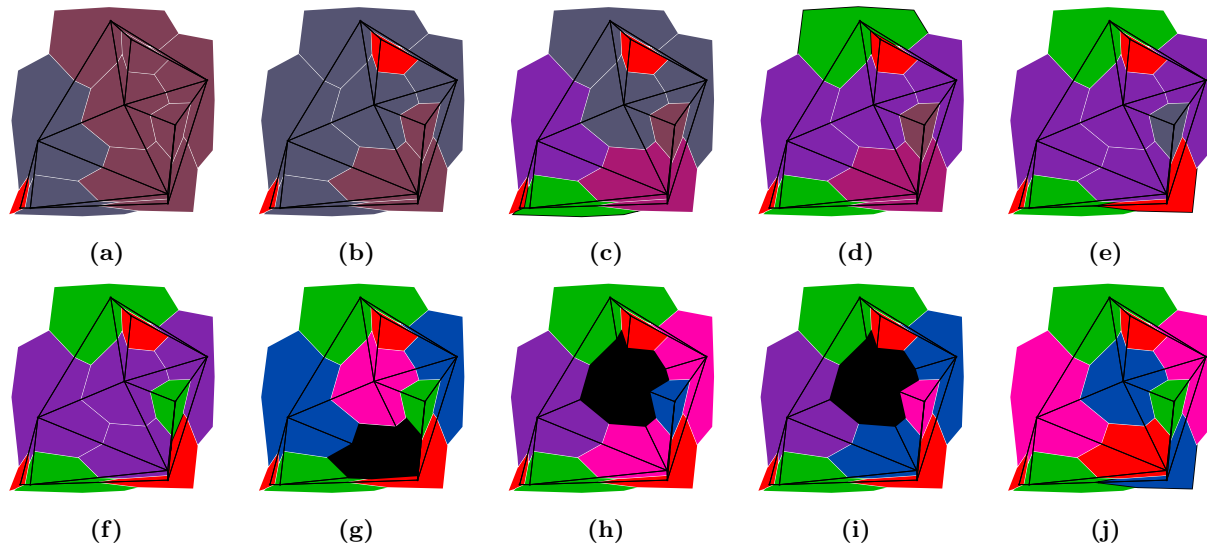
Figure 8

6 BACKTRACKING WITH CONSTRAINT PROPAGATION

6.1 Description

6.2 Experimental Approach

6.3 Results

7 GENETIC ALGORITHM

7.1 Implementation

7.2 Experimental Approach

7.3 Results

8 COMPARING ALGORITHM PERFORMANCE

8.1 Experimental Approach

8.2 Results

9 SUMMARY

REFERENCES

[1] Wikipedia, the free encyclopedia. Four color theorem, 2016. `https://upload.wikimedia.org/wikipedia/commons/4/4a/World_map_with_four_colours.svg`.

[2] Stuart Russel and Peter Norvig. *Artificial Intelligence: A Modern Approach*. Pearson Education, Upper Saddle River, New Jersey 07458, 3rd edition, 2010.

[3] Wikipedia, the free encyclopedia. Four color theorem, 2016. `https://en.wikipedia.org/wiki/Four_color_theorem#/media/File:Four_Colour_Planar_Graph.svg`.

[4] André LaMothe. *Tricks of the 3D Game Programming Gurus*. Sams Publishing, 201 West 103rd Street, Indianapolis, Indiana 46290, 2003.

[5] Thomas Cormen, Charles Leiserson, Ronald Rivest, and Clifford Stein. *Introduction to Algorithms*. Massachusetts Institute of Technology, Cambridge, Massachusetts, 3rd edition, 2009.

[6] Solomon W. Golomb and Leonard D. Baumert. Backtrack programming. *J. ACM*, 12(4):516–524, October 1965. `http://doi.acm.org/10.1145/321296.321300`.

[7] Robert M. Haralick and Gordon L. Elliott. Increasing tree search efficiency for constraint satisfaction problems. *Artificial Intelligence*, 14(3):263 – 313, 1980. `http://www.sciencedirect.com/science/article/pii/000437028090051X`.