

CSCI 446 Artificial Intelligence Project 1 Final Report

ROY SMART

NEVIN LEH

BRIAN MARSH

September 25, 2016

1 INTRODUCTION

The *Graph Coloring Problem* (GCP) is the problem of attempting to color a set of interconnected vertices, using a limited set of colors, such that no vertex has the same color as its neighbors. This problem is best visualized as the map coloring problem. As an example of the map coloring problem, consider the problem of assigning one of four colors to every country in the world, such that no two adjacent countries have the same color, as in Figure 1. It can be shown that the map coloring problem reduces to the graph coloring problem[2]

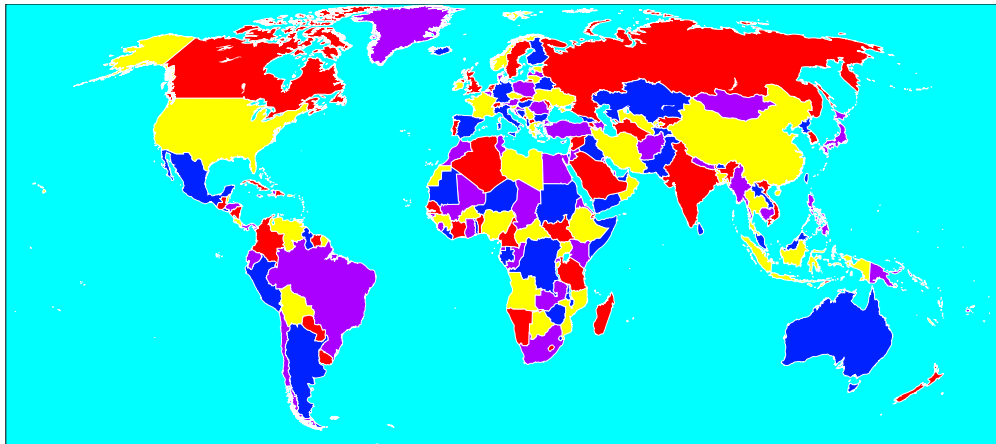


Figure 1: Map of the world satisfying the map coloring problem [1].

if we represent the countries as the vertices of the graph, and the borders between countries as the edges of the graph, shown in Figure 2. This configuration produces a *planar* graph, a graph with no edge intersections. For this project, we are asked to solve the GCP problem for planar graphs only. Since the problems are




Figure 2: Diagram describing how the map coloring problem can be transformed into the graph coloring problem [3].

guaranteed to be planar, we will represent our graphs using polygon maps generated by our program.

To solve the GCP we employed five different algorithms: Minimum Conflicts, Simple Backtracking, Backtracking with Forward Checking, Backtracking with Constraint Propagation (MAC), and Local Search

using a Genetic Algorithm. To evaluate these algorithms, we built a problem generating program that can produce a random set of planar graphs. Using the problem generator, we calculated a set of graphs between the sizes $\{10, 20, 30, \dots, 100\}$ and called the five graph coloring algorithms to solve the GCP. We measured GCP algorithm performance by measuring the number of read/write operations, the time required to find a solution, and the number of function calls for each algorithm. Using these metrics, we predict that the Minimum Conflicts algorithm will be the fastest, based off its performance on the eight-queens problem[2].

The code for this project is implemented in C++ and depends on the programs `cairo`, `gnuplot`, `ffmpeg`, and `LATEX` for graph output, performance output, video output, and documentation respectively.

2 PROBLEM GENERATION

We created a function that could produce many examples of planar graphs from a set of randomly scattered points. For an arbitrary set of points, there is no unique planar graph that can be constructed. To solve this issue, the problem statement has provided a algorithm for calculating an planar graph.

Select some point X at random and connect X by a straight line to the nearest point Y such that X is not already connected to Y and line crosses no other line. Repeat the previous step until no more connections are possible.

For our implementation of the above algorithm, we started by creating a complete graph (where each vertex is connected to every other vertex). Each vertex was associated with a list of edges, sorted by length. We then selected a point at random and inspected the first unchecked edge E . If E did not cross any of the accepted edges, it is added to the list of accepted edges and marked as checked. This process was repeated until every edge was marked as checked.

The most challenging part of the implementation described above was determining if two line segments intersected. An algorithm to determine if two edges (or line segments) cross has been outlined by LaMothe [4] and described in detail here. Let's start by defining two line segments

$$\begin{aligned} A &= \{\mathbf{a}, \mathbf{a}'\} \\ B &= \{\mathbf{b}, \mathbf{b}'\} \end{aligned} \tag{1}$$

where \mathbf{a} , \mathbf{a}' , \mathbf{b} , and \mathbf{b}' are vectors from the origin to the ends of the line segments. Next, compute the direction vectors for each line segment

$$\begin{aligned} \boldsymbol{\alpha} &= \mathbf{a}' - \mathbf{a} \\ \boldsymbol{\beta} &= \mathbf{b}' - \mathbf{b} \end{aligned} \tag{2}$$

Now, the trick to solving this problem is to parameterize the line segment using the direction vector and a parameter t_i in the domain $[0, 1]$.

$$\begin{aligned} \mathbf{p} &= \mathbf{a} + \boldsymbol{\alpha}t_a, \quad t_a \in [0, 1] \\ \mathbf{q} &= \mathbf{b} + \boldsymbol{\beta}t_b, \quad t_b \in [0, 1] \end{aligned} \tag{3}$$

From here, the solution is obvious. To find the point of intersection, we set $\mathbf{p} = \mathbf{q}$ and solve for the values of t_a and t_b . Then, if the solution is outside the domain of t_i , the line segments do not intersect. Thus,

$$\Rightarrow \begin{cases} p_x = q_x \\ p_y = q_y \end{cases} \tag{4}$$


$$\Rightarrow \begin{cases} a_x + \alpha_x t_a = b_x + \beta_x t_b \\ a_y + \alpha_y t_a = b_y + \beta_y t_b. \end{cases} \tag{5}$$

Equation 5 is a system of two equations and two unknowns, solving for t_a and t_b gives


$$\begin{aligned} t_a &= \frac{(-a_y + b_y)\beta_x + (a_x - b_x)\beta_y}{\alpha_y\beta_x - \alpha_x\beta_y} \\ t_b &= \frac{(a_y - b_y)\alpha_x + (-a_x + b_x)\alpha_y}{-\alpha_y\beta_x + \alpha_x\beta_y}. \end{aligned} \quad (6)$$

Finally, we can determine whether A and B intersect if $0 < t_a < 1$ and $0 < t_b < 1$ evaluates to true. This method may be more inefficient than other algorithms, such as the example outlined by Cormen[5] due to the division operation in Equation 6.

Using the expression above to determine line intersections, we were able to implement a problem generator that was reasonably efficient, generating approximately 1×10^3 examples per second. This efficiency was important, because it made testing and debugging the GCP algorithms much faster. Example output from the problem generator is shown in Figure 3a.



(a) Planar graph with 100 vertices.



(b) Polygon map corresponding to (a).

Figure 3: Visualization of the output of the problem generator. First a planar graph was created (a), and then a corresponding polygon map was found (b).

We have also implemented a way to produce a map of regions that is analogous to the graphs created by the problem generator, shown in Figure 3b. Since the graph coloring problem does not uniquely define the map coloring problem, we had to invent a metric that produces appropriate graphs. To define our metric, consider that if a polygon G_i corresponding to each point P_i were constructed out of the midpoints $M_{i,m}$ of the associated edges $E_{i,m}$ of P_i , then an adjacent polygon G_j would only touch at the corners, that is, our map would be full of gaps. Therefore, to fix the problem we also define the centroid $C_{i,n}$ of each triangle $T_{i,n}$ formed by the two adjacent edges $E_{i,m}$ and $E_{i,m+1}$. We then define the final polygon for P_i using the points $M_{i,m}$ and $C_{i,n}$. This procedure is unnecessary in terms of the assignment, but is helpful for debugging GCP algorithms, demonstrating example output, and for aesthetic pleasure.

3 MINIMUM CONFLICTS

4 SIMPLE BACKTRACKING

4.1 Description

Backtracking is a method of solving constraint satisfaction problems through an "educated" exhaustive search. It works by building a partial solution one component at a time and then testing the solution to check if the solution still has a chance of success [6]. If the solution cannot possibly succeed, the algorithm *backtracks* by moving back one component, modifying the component to form a new partial solution and then continuing. By checking partial solutions, backtracking can eliminate large portions of the solution-space, thus yielding greater efficiency than a simple exhaustive search.

In the graph coloring problem, backtracking assigns colors to one vertex at a time and checks to see if it has the same value as any of its neighbors. If none of the possible colors work for the current vertex, it backs up and tries a different color on a previously assigned vertex. If a color has no conflicts, backtracking continues to the next vertex in the sequence.




Figure 4: A map coloring example using simple backtracking on a graph with 10 vertices. From this diagram we will define a label for each vertex by the panel during which it is assigned, for example vertex (a) is colored red in panel 4a, vertex (c) is colored green in the panel 4c, etc. Note we will define vertex (j) as the last uncolored node, since panel 4j changes the value on vertex (h).

To illustrate the process of backtracking consider the sequence of images in Figure 4. This sequence follows the backtracking algorithm assigning the first acceptable color (in the order red, green, pink) to each vertex until the first backtracking event in panel 4j. The colors assigned to the vertices (a) through (g) seem reasonable until blue is assigned to vertex (h) in panel 4h. Then in panel 4i we can see that the only acceptable color for vertex (i) is pink and even worse, vertex (j) has cannot be assigned a color without being in conflict with its neighbors. Therefore, in panel 4j the algorithm is forced to backtrack to vertex (h) and try pink (which we will see).

The example above seeks to describe the major issue with backtracking: it can get stuck in situations where it can take an inordinate amount of time to backtrack far enough to fix the problem. In panel (j), after the backtracking operation, we can see that the algorithm did not retreat far enough, changing vertex (g) to pink still results in a conflict in vertex (j). The algorithm needs to back up to vertex (e) and change it to blue to resolve the conflict on vertex (j) and solve the graph. We will see in the following sections on forward checking and constraint propagation if we can mitigate this problem.

To implement simple backtracking we followed the recursive algorithm outlined in Russel and Norvig Section 6.3 [2]. Russel’s BACKTRACK loops through each value in the domain, and if the value is consistent with the constraints it recursively calls BACKTRACK. The function returns `true` when every variable has been assigned, and `false` if it cannot satisfy the constraints. The recursive function is advantageous, as it saves the state of the stack before each recursive call, making it trivial to restore the state of the solution after a failed recursive call to backtracking.

The crux of the backtracking procedure in the graph coloring problem is checking whether a coloring is consistent with constraints. From profiling our code, it was obvious that a fast constraint checking function was critical to the ability of backtracking to solve graph coloring problems up to $N = 100$ vertices in a reasonable amount of time. This facilitated a change of design from an object-oriented approach, to an array-based approach.

4.2 Experimental Approach

To evaluate backtracking’s performance on the graph coloring problem, we ran the algorithm on ten graph sizes $N = (10, 20, \dots, 100)$ with 20 trials for each value, for a total of 200 backtracking experiments. From each experiment we recorded the number of vertex reads, vertex writes, and number of recursive calls. Time constraints on this project demanded that we set a limit on the number of calls to backtracking, therefore we also recorded if this limit was reached for each experiment.

We expect that the number of recursive calls and the number of vertex writes and the number of vertex reads to be correlated by some constant factor. This is because there is a maximum of two writes and N reads for each recursive call. Therefore, we can use this information to verify that the algorithm is working properly.

4.3 Results




Figure 5: Logarithmic plot describing the performance of the backtracking algorithm vs. the number of vertices. The shaded regions represent the minimum and maximum values of each quantity.

5 BACKTRACKING WITH FORWARD CHECKING

5.1 Description

Backtracking with forward checking incorporates the usual simple backtracking algorithm described in the preceding section, while incorporating an ability to look into the future to see if a particular assignment would create an inconsistency. Forward checking achieves this by trying to make an inconsistency occur as fast as possible by checking if all current assignments are consistent with each other [7].

To accomplish forward checking in the graph coloring problem, we need to modify our representation of the graph to allow superpositions of colors and initialize the graph such that all colors are possible for each vertex. Next, for each call to backtracking, we assign a color to the next vertex, and delete that color from the vertex's neighbors list of possible colors. If any vertex has its list of possible colors completely eliminated, the algorithm backtracks and changes a color further up the tree.[2].




Figure 6: The map coloring example from Figure 4 using backtracking with forward checking. Note that the intermediate colors represent additive mixings of the four pure colors (red, green, blue, and pink), for example purple (the color of vertex (h) during panel (c)) is a combination of blue and pink, magenta is a combination of red and blue, etc.

To visualize how backtracking with forward checking is different from simple backtracking, we will consider again the map coloring example from the preceding section. To represent the superposition of colors required by this algorithm, we have formed a color map that represents additive mixings of the four pure colors. In this scheme, the color with all colors activated is mauve [RGBP]. The example starts with the graph initialized to mauve. Then vertex (a) is set to red, deleting red from its neighbors, vertices (h) and (c), forming gray [GBP]. After a similar process for vertex (b), vertex (c) is assigned green, deleting green from vertices (e) and (i) to form magenta [RBP] and from vertex (h) to form purple [RP].

This process produces similar results to simple backtracking until vertex (g) is colored blue in panel 6g, it then deletes blue from vertex (j), leaving pink. Next, in panel 6h, vertex (h) is chosen to be blue, leaving vertex (i) with pink as its only possible color. In panel 6i the algorithm selects the only choice for vertex (i) (pink), leaving vertex (j) with the null set. Since the null set is not allowed, the algorithm backtracks to vertex (g) in panel 6j, and changes its color to pink.

In this example, forward checking had no advantages over simple backtracking, as it discovered the inconsistency in vertex (j) in approximately the same number of steps. However, forward checking offers immense advantages for large graphs, where it can check if a color is consistent with its neighbors much sooner than simple backtracking, leading to much faster solutions.

For our implementation, we developed a function denoted `forward_check(vertex)` that checks the color of `vertex` and deletes that color from any neighboring vertices. If any neighboring vertex does not have any possible colors left, `forward_check()` returns `false` otherwise it returns `true`. This function is called by the same BACKTRACK function described in Section 4.1, with the additional instruction that `forward_check()` must evaluate to `true` before the next recursive call to BACKTRACK.

5.2 Experimental Approach

Our experimental approach for this algorithm is essentially the same as the approach for simple backtracking described in Section 4.2. We will still be measuring the number of vertex reads, vertex writes, and the number of recursive calls. However, since the forward checking algorithm must carry out more decisions for each recursive call, we expect there to be a larger amount of vertex reads and vertex writes for each recursive call. We will see in Section 8 if this additional overhead was worth it compared to simple backtracking. As with simple backtracking there is also a limit on how long each experiment can run, but this ends up not being important, as

5.3 Results




Figure 7: Logarithmic plot describing the performance of the backtracking algorithm vs. the number of vertices. The shaded regions represent the minimum and maximum values of each quantity.

6 BACKTRACKING WITH CONSTRAINT PROPAGATION

6.1 Description

Backtracking with constraint propagation uses the idea of *maintaining arc consistency* (MAC) to try to improve upon the performance of backtracking with forward checking. This variation of the backtracking algorithm can be described as recursive forward checking [2]. Where backtracking with forward checking assigns a color and deletes that color possibility from its neighbors, backtracking with constraint propagation assigns a color, deletes that color from its neighbors, and then uses that information to delete additional color possibilities from its neighbor's neighbors. This is the idea of constraint propagation. Furthermore, this

algorithm makes sure to maintain arc-consistency by making sure that every vertex retains a color that is consistent with the rest of the graph.




Figure 8: The map coloring example from Figure 4 using backtracking with constraint propagation.

To see how backtracking with constraint propagation improves upon the idea of backtracking with forward checking, we turn one last time to the example introduced in Section 4.1. Start by comparing Figures 6 and 8. Notice how panels up to 6f and 8f are precisely the same. Backtracking with forward checking and backtracking with constraint propagation haven't diverged up to this point since this is a small graph. The first change occurs when vertex (g) is assigned the color blue in panel 8g. In the forward checking example this reduces the color possibilities on vertex (j) to pink in panel 6g. But in the constraint propagation example it propagates this information to vertex (h) and (i) in panel 8g, reducing them to blue and null, respectively. Since constraint propagation noticed this error two steps earlier than forward checking, it is able to quickly recurse back in panels 8h and 8i far enough to change vertex (e) to blue in panel 8j. This change allows the solution to be found almost immediately, as the next choice for vertex (f) is green, which constrains vertex (g) to pink, (j) to blue, (h) to pink, and (i) to red.

While the example above may seem only incrementally better than forward checking, as we will see in Section 8, constraint propagation with MAC adds another large improvement to the backtracking algorithm from forward checking. This is because this algorithm can prune off large sections of the search tree by performing a search of the local neighborhood at each recursive backtracking step.

In our implementation of backtracking with constraint propagation we used the algorithm AC-3 introduced by Mackworth in 1977 [8]. This algorithm is described in depth by Figure 6.3 of Russel and Norvig [2]. At the heart of AC-3 is a function denoted REVISE. The task of REVISE is to loop through the possible colors i in one vertex and make sure that it is consistent with all the possible colors j in a neighboring vertex. If any i is incompatible with all of j , it is deleted from the list of possible colors, and REVISE returns **true**, indicating that the domain i has been reduced. Otherwise, REVISE returns **false** since i remained unchanged.

AC-3 uses a queue of remaining vertices and the function REVISE to propagate constraints across the whole graph, while simultaneously maintaining arc consistency. For each vertex in the queue AC-3 calls REVISE to see if the vertex provided any new information. If REVISE returns **true**, AC-3 first checks if the domain i is zero, if so, AC-3 returns **false**, signifying an arc-consistency violation. If the domain of i is nonzero, AC-3 propagates the information obtained through REVISE by adding the vertex's neighbors to the queue. AC-3 continues until its queue is empty and returns **true**.

Similarly to backtracking with forward checking, backtracking with constraint propagation makes a small modification to the function BACKTRACK, but instead of calling `forward_check()` before the next recursive call, we will ensure that AC-3 evaluates to **true** before the next recursive call to BACKTRACK.

6.2 Experimental Approach

Again, the experimental approach for this backtracking algorithm does not differ greatly from the previous two versions of backtracking. We do note however, that constraint propagation and MAC do have an increased read and write overhead as compared to forward checking and simple backtracking. Since constraint propagation introduces even more overhead than forward checking on top of backtracking, in Section 8, we will see at what point, if ever, constraint propagation outpaces backtracking.

6.3 Results




Figure 9: Logarithmic plot describing the performance of the backtracking with constraint propagation algorithm vs. the number of vertices. The shaded regions represent the minimum and maximum values of each quantity.

7 GENETIC ALGORITHM

7.1 Implementation

7.2 Experimental Approach

7.3 Results

8 COMPARING ALGORITHM PERFORMANCE

8.1 Experimental Approach

As described in the preceding sections, the number of vertex reads and number of vertex writes were our main metrics for comparing algorithm performance. Under these metrics, all calculations performed by each algorithm were assumed to take the same amount of time between subsequent reads and writes. This assumption is obviously untrue, but tests on our algorithm showed

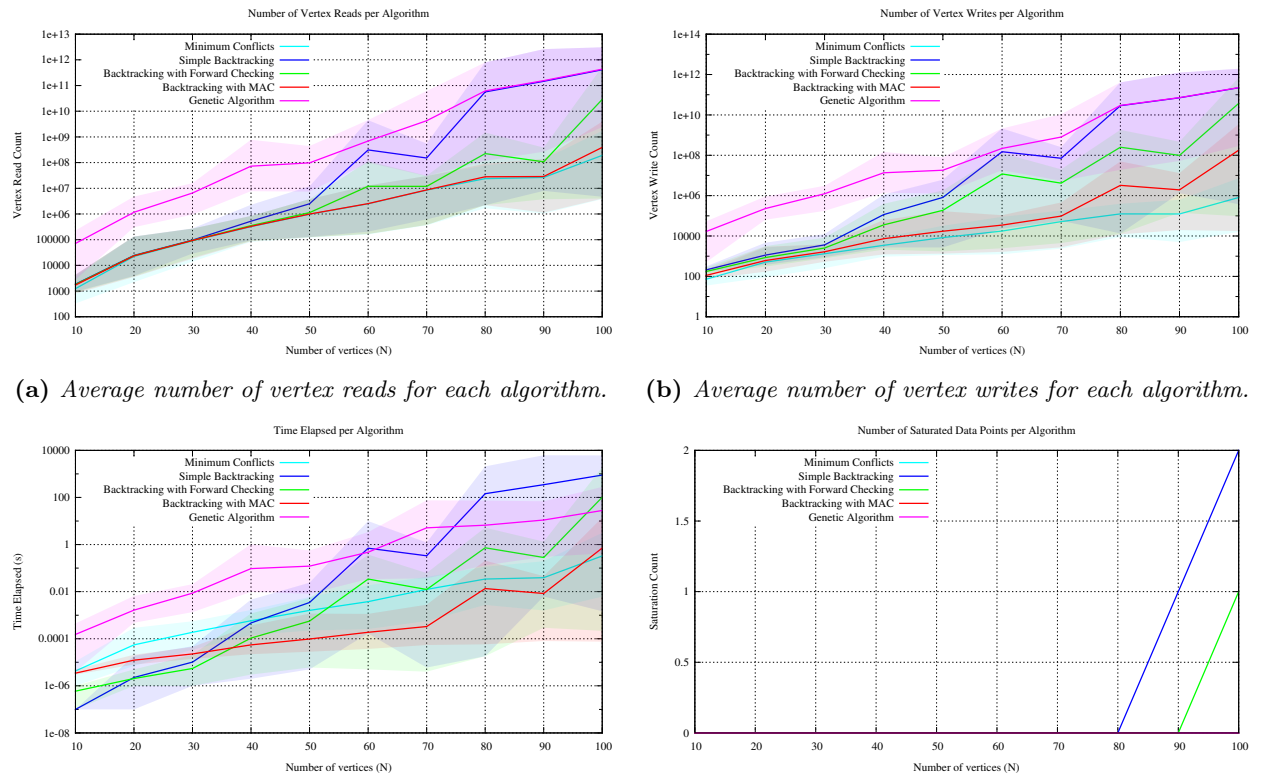


Figure 10: Logarithmic plots comparing the performance of the five different GCP algorithms against four different measurements. Shaded regions represent the maximum and minimum values.

8.2 Results

9 SUMMARY

REFERENCES

- [1] Wikipedia, the free encyclopedia. Four color theorem, 2016. https://upload.wikimedia.org/wikipedia/commons/4/4a/World_map_with_four_colours.svg.
- [2] Stuart Russel and Peter Norvig. *Artificial Intelligence: A Modern Approach*. Pearson Education, Upper Saddle River, New Jersey 07458, 3rd edition, 2010.
- [3] Wikipedia, the free encyclopedia. Four color theorem, 2016. https://en.wikipedia.org/wiki/Four_color_theorem#/media/File:Four_Colour_Planar_Graph.svg.
- [4] André LaMothe. *Tricks of the 3D Game Programming Gurus*. Sams Publishing, 201 West 103rd Street, Indianapolis, Indiana 46290, 2003.
- [5] Thomas Cormen, Charles Leiserson, Ronald Rivest, and Clifford Stein. *Introduction to Algorithms*. Massachusetts Institute of Technology, Cambridge, Massachusetts, 3rd edition, 2009.
- [6] Solomon W. Golomb and Leonard D. Baumert. Backtrack programming. *J. ACM*, 12(4):516–524, October 1965. <http://doi.acm.org/10.1145/321296.321300>.
- [7] Robert M. Haralick and Gordon L. Elliott. Increasing tree search efficiency for constraint satisfaction problems. *Artificial Intelligence*, 14(3):263 – 313, 1980. <http://www.sciencedirect.com/science/article/pii/000437028090051X>.
- [8] Alan K. Mackworth. Consistency in networks of relations. *Artificial Intelligence*, 8(1):99 – 118, 1977. <http://www.sciencedirect.com/science/article/pii/0004370277900078>.